



Conducted By Kalyan Citizens' Education Society

B. K. Birla College of Arts, Science & Commerce, Kalyan

(Empowered Autonomous Status)

Affiliated to University of Mumbai

ISO 9001: 2015 Certified | Reaccredited by NAAC (4th Cycle) with 'A++' Grade (CGPA – 3.51)

B. K. Birla College Campus Road, Kalyan (West) - 421301

CERTIFICATE

This is to certify that **Mr./Ms.** _____

Roll Number/Seat Number _____ a student of
Programme _____ has successfully completed
the practical/project of the **Course** _____
for **Semester** _____ under the valuable guidance of
_____ during the **A.Y.** _____.

Date:

Place:

Internal Examiner

External Examiner

Seal

INDEX

Programme: _____ **Semester:** _____

Course: _____

Sr. No.	Prac. No.	Particulars	Date	Teacher Signature
1.	1.			
2.	2.			
3.	3.			
4.	4.			
5.	5.			
6.	6.			
7.	7.			
8.	8.			

Case Study 1: AI-Powered Chatbots on Cloud for Customer Support Automation

Objective:

To analyze the architecture, implementation, and business impact of deploying AI-powered chatbots on cloud platforms. This case study explores how services like **AWS Lex**, **Google Dialogflow**, and **Azure Bot Service** leverage Natural Language Processing (NLP) and serverless computing to create scalable and intelligent customer support solutions.

Background:

In today's digital-first landscape, customers expect immediate, 24/7 support. Traditional support models, relying on human agents, are often costly and struggle to scale during peak demand. AI-powered chatbots address this challenge by automating responses to frequently asked questions and handling routine tasks. By hosting these bots on the cloud, companies can offload complex infrastructure management, benefit from pay-as-you-go pricing, and integrate with a vast ecosystem of cloud services.

Technical Architecture Deep Dive:

A modern cloud chatbot architecture is a decoupled, event-driven system composed of several managed services.

1. **Frontend/Client Interface:** This is the user-facing component. It can be a web chat widget (using JavaScript), a mobile application (iOS/Android), or an integration with third-party platforms like Slack, WhatsApp, or Facebook Messenger.
2. **API Gateway:** Services like **Amazon API Gateway** or **Azure API Management** act as the front door for all incoming user requests. It provides security, throttling, and routes requests to the appropriate backend service.
3. **NLP/NLU Engine (The "Brain"):** This is the core AI service that processes the user's text.
 - o **AWS Lex, Google Dialogflow, or Azure LUIS (Language Understanding)** are used here.
 - o **Natural Language Processing (NLP):** The service first processes the raw text.
 - o **Natural Language Understanding (NLU):** It then performs two key tasks:
 - **Intent Recognition:** Identifies the user's goal (e.g., CheckOrderStatus, ResetPassword).
 - **Entity Extraction:** Pulls out key pieces of information (e.g., order_id: "12345", username: "john.doe").
4. **Business Logic Layer (Serverless Functions):** Once the intent and entities are identified, the NLP service triggers a backend function to execute the required action.
 - o **AWS Lambda, Google Cloud Functions, or Azure Functions** are perfect for this.
 - o This function contains the custom code to, for instance, query a database, call an external API, or perform a calculation.
5. **Data Layer (Cloud Database):** The backend function often needs to interact with a database to retrieve or store information.
 - o **NoSQL Databases** like **Amazon DynamoDB** or **Google Firestore** are commonly used for their high speed and scalability, which is ideal for storing conversation logs or user data.
 - o **Relational Databases** like **Amazon RDS** or **Azure SQL Database** might be used to access existing business data, such as customer order information.

6. **Response Generation:** The backend function formats a response, which is sent back through the API Gateway to the user on the frontend.

Workflow Example: Checking an Order Status:

User: "What is the status of my order number 98765?"

1. **Frontend:** The chat UI sends the user's message as a JSON payload to the API Gateway endpoint.
2. **API Gateway:** Authenticates the request and forwards it to the NLP Engine (e.g., Dialogflow).
3. **NLP Engine (Dialogflow):**
 - Recognizes the **Intent** as CheckOrderStatus.
 - Extracts the **Entity** orderNumber with the value "98765".
 - Triggers the associated serverless function (e.g., a Google Cloud Function).
4. **Logic Layer (Cloud Function):**
 - The function receives the intent (CheckOrderStatus) and the entity (orderNumber: "98765").
 - It executes code to query the company's Orders Database (e.g., an RDS instance) using the provided order number.
 - The database returns the status: "Shipped".
5. **Response Generation:** The function constructs a user-friendly response string: "Your order #98765 has been shipped and is on its way!"
6. **Return Path:** The response is sent back through the API Gateway to the user's chat window.

Advantages

- **24/7 Availability:** Provides instant, round-the-clock automated support without human intervention.
- **Hyper-Scalability:** Leverages serverless functions and managed services that automatically scale from zero to thousands of requests per second.
- **Cost Reduction:** Dramatically reduces the cost per interaction compared to human agents, freeing them up for more complex issues.
- **Data-Driven Insights:** Captures and analyzes user queries to identify common pain points, popular products, or gaps in service.
- **Omnichannel Integration:** A single chatbot backend can be easily deployed across multiple platforms (web, mobile, social media, voice assistants).
- **Consistent Service:** Ensures every customer receives the same quality and accuracy of information, aligning with brand guidelines.

Challenges and Considerations

- **Handling Ambiguity:** Chatbots can struggle with complex, ambiguous, or slang-filled user queries. A robust fallback mechanism to escalate to a human agent is crucial.
- **State Management:** Maintaining conversation context (remembering what was said earlier) can be complex and requires careful design.
- **Security:** As chatbots often handle sensitive user data, ensuring secure authentication, authorization, and data encryption is paramount.
- **Maintenance Overhead:** Intents, entities, and response scripts must be continuously updated and trained as business processes and customer queries evolve.

Example in Practice: Domino's Pizza

The Domino's Pizza chatbot ("Dom") is a prime example of cloud AI in action. Built with **Google Dialogflow**, it allows customers to order pizzas through natural conversation via text, social media, or voice.

- **Intent:** OrderPizza
- **Entities:** pizzaType (Pepperoni), size (Large), crust (Thin Crust), and address (User's saved address).
- **Backend Logic:** The Dialogflow agent integrates with Domino's backend ordering system via APIs to place the order, process payment, and provide a real-time delivery tracker link.

Conclusion

Cloud-based AI chatbots represent a fundamental shift in customer service, moving from a reactive, human-dependent model to a proactive, automated one. By leveraging the power of managed NLP services and serverless computing, organizations can build intelligent, scalable, and cost-effective solutions that significantly enhance customer experience and drive operational efficiency.

Case Study 2: AI Document Analysis using Cloud OCR (Invoice Extraction)

Objective

To investigate how cloud-based AI services for Optical Character Recognition (OCR) and document analysis, such as **AWS Textract**, **Azure Form Recognizer**, and **Google Document AI**, can automate the extraction of structured data from documents like invoices, receipts, and forms.

Background

Finance, logistics, healthcare, and legal departments in most organizations are inundated with vast quantities of semi-structured documents. The manual process of reading these documents and keying data into systems like ERPs or databases is slow, expensive, and prone to human error. Intelligent Document Processing (IDP) systems built on cloud AI services provide a scalable and accurate solution, transforming unstructured document images into structured, machine-readable data.

Technical Architecture Deep Dive

This process is typically implemented as an event-driven, serverless pipeline on the cloud, ensuring scalability and cost-efficiency.

1. **Ingestion Point (Cloud Storage):** Documents (e.g., PDF, JPG, PNG) are uploaded to a secure, scalable object storage service.
 - **Services:** **Amazon S3**, **Azure Blob Storage**, **Google Cloud Storage**.
 - An upload event (e.g., S3:ObjectCreated:*) is configured to automatically trigger the next step in the pipeline.
2. **Processing Trigger (Serverless Function):** A lightweight, serverless function is initiated by the storage event. Its role is to orchestrate the analysis.
 - **Services:** **AWS Lambda**, **Azure Functions**, **Google Cloud Functions**.
 - This function receives the document's location (e.g., bucket name and file key) as input.
3. **Intelligent OCR & Data Extraction (The "Brain"):** The serverless function calls the core AI document analysis service. This is more advanced than traditional OCR because it understands document structure.
 - **Services:** **AWS Textract**, **Azure Form Recognizer**, **Google Document AI**.
 - **Functionality:**
 - **Raw Text OCR:** Extracts all plain text from the document.
 - **Forms Extraction:** Identifies and extracts data from key-value pairs (e.g., "Invoice Number": "INV-1234").
 - **Table Extraction:** Recognizes tabular data and preserves the row and column structure.
 - **Handwriting Recognition:** Can often process handwritten text in addition to printed text.
 - The service returns a structured JSON object containing the extracted data, confidence scores, and bounding box coordinates for each piece of information.

4. **Post-Processing & Validation (Logic Layer):** The same serverless function (or a subsequent one) receives the JSON output from the AI service.
 - **Parsing:** It parses the JSON to isolate critical fields (e.g., vendor name, total amount, due date).
 - **Validation:** It can perform business rule checks, such as verifying the date format or ensuring the sum of line items matches the total amount.
 - **Enrichment:** It might enrich the data by cross-referencing the vendor name with an existing database to get a vendor ID.
5. **Data Persistence (Cloud Database):** Once validated, the clean, structured data is stored in a suitable database.
 - **NoSQL Databases (e.g., Amazon DynamoDB):** Ideal for storing the flexible JSON output for quick retrieval.
 - **Relational Databases (e.g., Amazon Aurora/RDS):** Used when the data needs to be integrated into existing structured systems like an ERP or accounting platform.
6. **Visualization and Integration (Downstream Systems):** The structured data is now ready for use.
 - **Dashboards:** Tools like **Power BI**, **Tableau**, or **Amazon QuickSight** can connect to the database to visualize spending trends, payment cycles, etc.
 - **ERP/Billing Systems:** The data can be programmatically pushed into systems like SAP, Oracle, or QuickBooks to automate the accounts payable process.

Workflow Example: Processing a Supplier Invoice

An accounts payable clerk uploads a supplier's PDF invoice into a designated Amazon S3 bucket.

1. **Ingestion:** The PDF file lands in the invoices-pending S3 bucket.
2. **Trigger:** An S3 PutObject event automatically invokes an **AWS Lambda** function.
3. **Extraction:**
 - The Lambda function calls the **AWS Textract** StartDocumentAnalysis API, pointing to the location of the PDF in S3.
 - Textract processes the document, identifying the invoice number, line items in a table, and the final amount due.
 - It returns a detailed JSON object with all this structured information.
4. **Post-Processing:**
 - The Lambda function parses the JSON, extracting the values for vendor_name, invoice_id, due_date, and total_amount.
 - It checks if total_amount is a valid number and if due_date is a valid date.
5. **Persistence:** The function then writes a new item to a **DynamoDB** table called Invoices, containing the clean, validated data.
6. **Integration:** A separate process (or another event trigger on the DynamoDB table) picks up the new record and automatically creates a draft bill in the company's accounting software via an API call.

Advantages

- **Eliminates Manual Data Entry:** Drastically reduces the time, cost, and errors associated with manual keying.
- **Accelerated Business Cycles:** Speeds up processes like accounts payable, claims processing, and customer onboarding from days to minutes.
- **High Accuracy & Confidence:** AI models provide confidence scores for each extracted field, allowing for a human-in-the-loop review only for low-confidence results.
- **Scalability:** The serverless architecture effortlessly scales to process thousands of documents per hour during peak periods without any infrastructure management.
- **Creates Searchable Archives:** Digitized documents and their metadata can be easily indexed, making them fully searchable for compliance and auditing purposes.

Challenges and Considerations

- **Template Variation:** The models work best with consistent layouts. Performance can degrade if invoices from different vendors have wildly different formats. Custom models may need to be trained for specific, complex documents.
- **Image Quality:** The accuracy of OCR is highly dependent on the quality of the source document (e.g., resolution, lighting, skew, handwritten notes).
- **Cost Management:** While cost-effective, processing millions of pages can become expensive. It's important to monitor API usage and optimize the workflow.
- **Data Privacy:** Documents may contain sensitive PII (Personally Identifiable Information). It's crucial to implement proper security controls, encryption, and access policies within the cloud environment.

Conclusion

Cloud-based Intelligent Document Processing is a transformative technology that allows organizations to unlock the value trapped in their unstructured documents. By building automated pipelines with services like AWS Textract or Azure Form Recognizer, companies can achieve significant gains in operational efficiency, data accuracy, and business agility, ultimately freeing up human capital to focus on higher-value tasks.

Case Study 3: Real-Time Cloud Analytics for Smart Cities

Objective

To examine the end-to-end cloud architecture required to build a real-time analytics platform for smart city management. This case study focuses on how the integration of IoT (Internet of Things) devices, streaming data services, and AI/ML models on the cloud enables proactive urban management and data-driven decision-making.

Background

Urbanization is leading to increasingly complex challenges in traffic management, public safety, resource allocation, and environmental sustainability. Smart cities leverage a network of IoT sensors—embedded in traffic lights, public utilities, waste bins, and environmental monitors—to collect vast amounts of real-time data. By harnessing cloud AI platforms, city planners and operators can move from a reactive to a predictive model, anticipating issues like traffic bottlenecks or air quality degradation before they become critical.

Technical Architecture Deep Dive

A smart city analytics platform is a multi-layered, highly available system designed to handle high-velocity data streams.

1. **Sensing & Edge Layer:** This layer consists of the physical IoT devices distributed across the city.
 - **Devices:** GPS sensors on buses, video cameras for traffic monitoring, air quality sensors, smart water meters, and smart lighting controllers.
 - **Edge Gateways:** These devices aggregate data from multiple sensors locally, perform initial filtering or compression, and securely transmit it to the cloud using efficient protocols like **MQTT** or **HTTP/S**.
2. **Ingestion & Messaging Layer:** This is the secure entry point for all IoT data into the cloud.
 - **Services:** **AWS IoT Core**, **Azure IoT Hub**, or **Google Cloud IoT Core**.
 - **Functionality:** These managed services can handle millions of device connections, provide device authentication and authorization, and offer a rules engine to route incoming data to different cloud services.
3. **Stream Processing Layer (The "Hot Path"):** Data that requires immediate analysis flows through this real-time pipeline.
 - **Services:** **Amazon Kinesis**, **Azure Stream Analytics**, or **Google Cloud Dataflow**.
 - **Functionality:** These services can perform real-time data transformation, aggregation, and filtering on the fly. For example, they can calculate the average traffic speed over a 1-minute window or detect an anomalous spike in pollution levels.
4. **AI / Machine Learning Layer:** The processed real-time data is fed into machine learning models for predictive insights.
 - **Services:** **Amazon SageMaker**, **Azure Machine Learning**, or **Google AI Platform**.
 - **Model Types:**
 - **Forecasting:** Time-series models to predict traffic congestion 30 minutes in advance.

- **Anomaly Detection:** Identifying unusual patterns, like a sudden drop in water pressure indicating a pipe leak.
- **Classification:** Analyzing camera feeds to classify available vs. occupied parking spots.

5. Data Storage Layer: The platform utilizes a dual-approach to data storage.

- **Hot Storage (Real-Time):** Data for immediate dashboarding is stored in a high-speed, low-latency database. Examples include **Amazon Timestream** (a time-series database) or **Redis**.
- **Cold Storage (Long-Term):** All raw and processed data is archived in a cost-effective data lake for historical analysis and training future ML models. Services include **Amazon S3**, **Azure Data Lake Storage**, or **Google Cloud Storage**.

6. Presentation & Action Layer: This is where insights are consumed and acted upon.

- **Dashboards:** Real-time dashboards built with **Amazon QuickSight**, **Microsoft Power BI**, or **Grafana** display key metrics for city operators.
- **Alerting:** Automated alerts are triggered based on predefined rules (e.g., air quality index > 150). Services like **Amazon Simple Notification Service (SNS)** can send notifications via email, SMS, or to a mobile app.
- **Actuation:** In advanced systems, insights can trigger automated actions, such as dynamically changing traffic light timings to ease congestion.

Workflow Example: Dynamic Traffic Management

A traffic accident occurs on a major city highway during rush hour.

1. **Sensing:** Roadside cameras and embedded magnetic loop sensors detect a sudden, complete stop in traffic flow in one lane.
2. **Ingestion:** The data (video snippets and sensor readings showing zero velocity) is sent via an edge gateway to **AWS IoT Core**.
3. **Processing:** IoT Core's rules engine forwards the data to **Amazon Kinesis**, which identifies this pattern as a high-priority anomaly.
4. **AI Analysis:** The anomaly data is fed to a pre-trained **SageMaker** model, which confirms the event is an accident (not just a traffic jam) and predicts a 90% probability of severe congestion spreading to adjacent routes within 15 minutes.

5. Action & Presentation:

- An automated alert is sent via **Amazon SNS** to the traffic management center and emergency services.
- The operator's real-time dashboard in **QuickSight** immediately shows the affected road segment in red.
- The system automatically triggers a command to update electronic road signs on feeder roads, suggesting alternative routes to drivers.

Advantages

- **Proactive Decision-Making:** Enables authorities to anticipate and mitigate problems, rather than just reacting to them.
- **Optimized Resource Allocation:** Data-driven insights help optimize public transport schedules, waste collection routes, and energy distribution.
- **Enhanced Public Safety:** Faster emergency response times through real-time incident detection and improved infrastructure monitoring.
- **Environmental Sustainability:** Helps manage pollution levels, conserve water, and reduce the city's carbon footprint through smart energy grids.
- **Scalability and Reliability:** The cloud infrastructure can seamlessly scale to accommodate an increasing number of sensors and data loads.

Challenges and Considerations

- **Data Privacy & Security:** Collecting city-wide data raises significant privacy concerns. Strong anonymization, encryption, and strict access control policies are essential.
- **High Initial Investment:** Deploying thousands of sensors and building the cloud infrastructure requires significant upfront investment.
- **Data Interoperability:** Integrating data from a wide variety of sensors and systems from different manufacturers can be a major technical hurdle.
- **Network Reliability:** The entire system depends on a reliable communication network (e.g., 5G, LoRaWAN) to transmit data from sensors to the cloud.

Conclusion

Cloud AI is the engine that powers the modern smart city. By providing a scalable and intelligent platform for IoT data analytics, it transforms raw sensor readings into actionable intelligence. This allows municipal governments to create safer, more efficient, and more sustainable urban environments for their citizens, fundamentally changing the way cities are managed.

Case Study 4: AI-Based Healthcare Diagnostics on Cloud

Objective

To analyze the architecture and impact of cloud-hosted artificial intelligence models in medical diagnostics. This case study explores how cloud platforms facilitate the entire lifecycle—from data ingestion and model training to secure clinical deployment—of AI tools that assist healthcare professionals in detecting diseases from medical imagery.

Background

Radiology and pathology are cornerstone diagnostic fields that rely on the interpretation of complex visual data (X-rays, CT scans, MRIs, tissue slides). The global shortage of specialists, combined with an increasing volume of scans, creates significant workload pressure and can lead to diagnostic delays or errors. Cloud AI offers a powerful solution by providing a secure, scalable infrastructure to train and deploy deep learning models. These models act as an assistive tool for clinicians, helping to prioritize critical cases, improve detection accuracy, and provide quantitative analysis. Critically, these systems must operate within strict regulatory frameworks like **HIPAA** (in the US) to ensure patient data privacy and security.

Technical Architecture Deep Dive

Building a clinical-grade AI diagnostic tool involves a compliant and robust cloud architecture.

1. Secure Data Ingestion & Anonymization:

- **Source:** Medical images, typically in **DICOM** format, are exported from a hospital's **PACS** (Picture Archiving and Communication System) or **EMR** (Electronic Medical Record) system.
- **Anonymization:** Before being uploaded, a crucial preprocessing step is performed to strip all Protected Health Information (PHI) from the DICOM metadata to comply with privacy regulations.
- **Cloud Storage:** The anonymized images are uploaded to a secure, compliant object storage service like **Amazon S3** (using S3 Object Lock for immutability) or **Azure Blob Storage** with HIPAA-compliant configurations.

2. Data Lake & Labeling:

- The anonymized images are stored in a centralized data lake. This repository becomes the single source of truth for model training.
- **Data Labeling:** To train a supervised learning model, a dataset must be labeled by medical experts. Services like **Amazon SageMaker Ground Truth** or **Google Cloud's Data Labeling Service** provide workflows where radiologists can draw bounding boxes around tumors or classify images (e.g., 'pneumonia present' vs. 'normal').

3. Model Training, Tuning, & Validation (The "Brain"):

- **Model Choice:** **Convolutional Neural Networks (CNNs)** are the state-of-the-art models for image analysis tasks like classification, object detection, and segmentation.
- **Training Platform:** Managed services like **Amazon SageMaker**, **Azure Machine Learning**, or **Google's Vertex AI** are used to train these complex models on powerful GPU instances. These platforms handle the underlying infrastructure, allowing data scientists to focus on the model itself.

- **Validation:** The model is rigorously tested against a separate, unseen dataset to validate its accuracy, sensitivity, and specificity before it can be considered for deployment.

4. Secure Deployment & Inference:

- **API Endpoint:** Once validated, the trained model is deployed as a highly available, scalable, and secure API endpoint using services like **Amazon SageMaker Endpoints** or a containerized application on **Azure Kubernetes Service (AKS)**.
- **Inference:** For a new patient case, the anonymized medical image is sent to this endpoint via a secure API call. The model processes the image and returns its prediction (e.g., a probability score of malignancy) in a structured format like JSON, usually within seconds.

5. Clinical Integration & Presentation Layer:

- The model's output is not a final diagnosis but a piece of decision-support information for the clinician.
- **Dashboard:** The results, sometimes including a "heatmap" that visually highlights suspicious regions on the image, are displayed on a secure web-based dashboard for the doctor to review.
- **PACS/EMR Integration:** In advanced setups, the AI results are sent back and integrated directly into the hospital's native radiology software (PACS), appearing as a secondary analysis alongside the original image.

Workflow Example: AI-Assisted Lung Nodule Detection

A radiologist is reviewing a patient's chest CT scan for signs of lung cancer.

1. **Image Upload:** The anonymized CT scan is securely uploaded from the hospital's PACS to a designated S3 bucket.
2. **Trigger & Inference:** The upload triggers an **AWS Lambda** function, which calls the deployed **SageMaker** model endpoint, passing the location of the new scan.
3. **AI Analysis:** The CNN model analyzes hundreds of images in the CT scan series and detects a small, suspicious nodule (8mm) that could be an early-stage tumor.
4. **Result Generation:** The model returns a JSON response containing:
 - The probability of malignancy (e.g., 0.85).
 - The coordinates and slice number of the detected nodule.
 - A link to a generated heatmap image.

5. Clinical Review:

- The case is automatically flagged and moved to the top of the radiologist's worklist.
- When the radiologist opens the case in their viewer, they see the original CT scan alongside the AI's findings: the highlighted nodule and the 85% malignancy score.
- The radiologist uses this information to pay closer attention to the flagged area, ultimately confirming the finding and recommending a follow-up biopsy. The AI acted as a vigilant "second reader."

Advantages

- **Improved Diagnostic Accuracy & Speed:** Acts as a second pair of eyes for clinicians, reducing the risk of missed findings and accelerating the review process.
- **Early Disease Detection:** AI models can often detect subtle patterns that are imperceptible to the human eye, enabling earlier diagnosis and better patient outcomes.
- **Workload Triage:** Automatically prioritizes critical cases (e.g., brain hemorrhage, pulmonary embolism), ensuring that the most urgent patients receive immediate attention.
- **Democratization of Expertise:** Makes expert-level analysis available to smaller clinics or in remote areas that may not have on-site specialists.
- **Scalable and Secure:** Cloud platforms provide the necessary compute power for training large models and offer robust security and compliance (e.g., HIPAA, GDPR) for handling sensitive medical data.

Challenges and Considerations

- **Regulatory Approval:** AI diagnostic tools are considered medical devices and require stringent validation and approval from regulatory bodies like the **FDA** in the US or an **CE mark** in Europe.
- **Data Bias and Generalizability:** A model trained primarily on data from one demographic may perform poorly on others. Ensuring diverse and representative training data is a major ethical and technical challenge.
- **The "Black Box" Problem:** It can be difficult to understand exactly *why* a deep learning model made a particular prediction, which can be a barrier to clinical trust and adoption.
- **Integration with Hospital IT:** Integrating the AI solution seamlessly into existing and often legacy hospital IT systems (like PACS and EMRs) is complex.

Conclusion

Cloud-based AI is poised to revolutionize medical diagnostics by empowering clinicians with powerful, data-driven insights. While significant challenges around regulation, ethics, and integration remain, these tools have the potential to enhance diagnostic accuracy, improve efficiency, and ultimately make high-quality healthcare more accessible and effective on a global scale.

Practical 1: Building a Cloud Chatbot using Amazon Lex

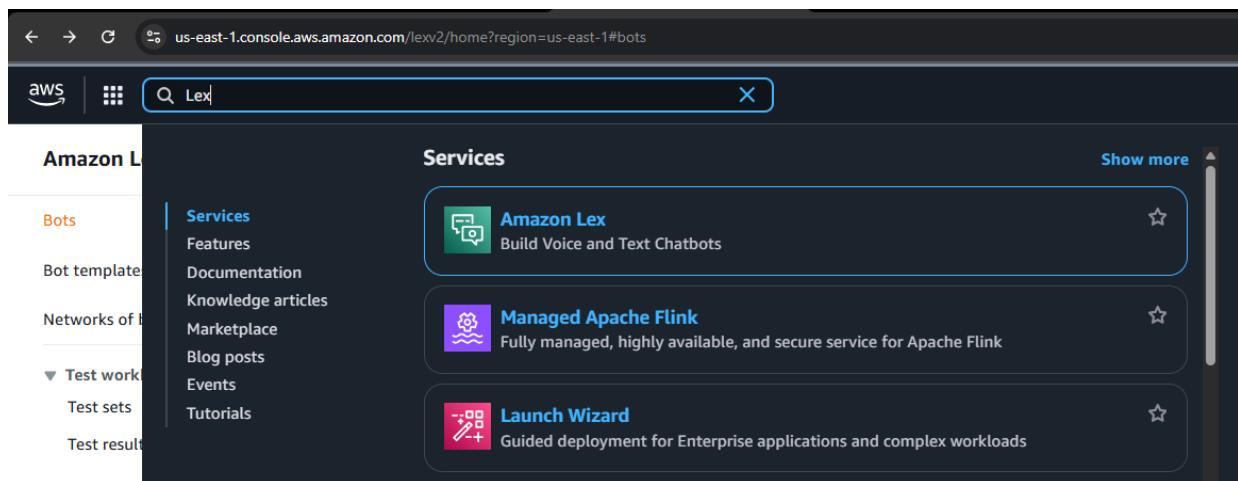
Aim: To create and test a basic conversational chatbot using Amazon Lex to handle simple, automated customer support queries.

Tools Used:

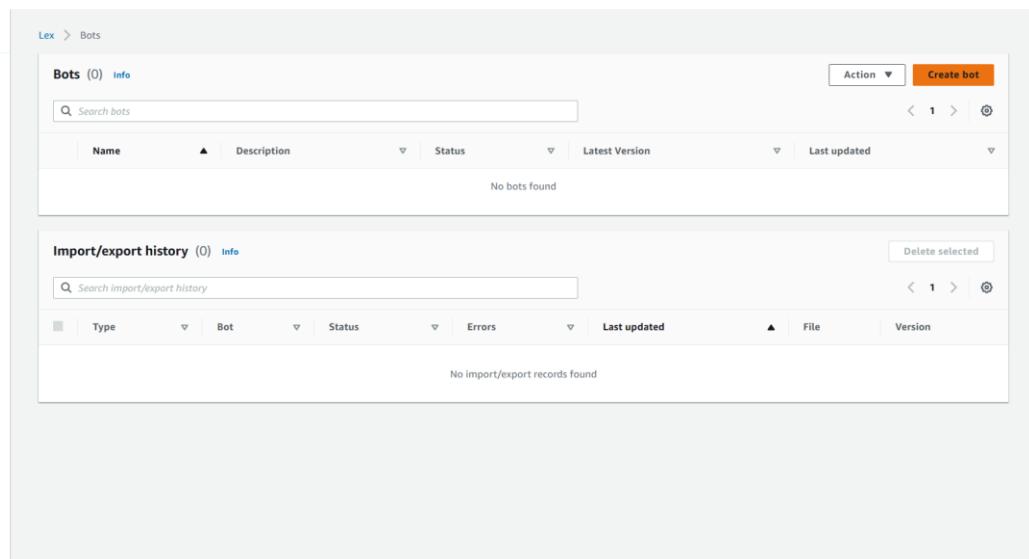
- AWS Management Console
- **Amazon Lex**
- **AWS Lambda** (Optional, for advanced logic)

Procedure:

1. Sign in to the AWS Management Console and navigate to the **Amazon Lex** service.



2. Click "Create bot" and choose the "Create a blank bot" option.



3. Name your bot (e.g., CustomerSupportBot) and configure the required IAM permissions.

AWS | Lex | Lex | X |

Lex > Bots > Create bot

Step 1
Configure bot settings

Step 2
Add languages

Configure bot settings Info

Creation method

Create a blank bot
Create a basic bot with no preconfigured languages, intents, and slot types.

Start with an example
An example bot has preconfigured languages, intents, and slot types. You can change these settings.

Start with transcripts
Automatically generate intents from conversation transcripts that you upload. Only English (US) language is available when starting with a transcript.

Bot configuration

Bot name

Maximum 100 characters. Valid characters: A-Z, a-z, 0-9, -, _

Description - optional
This description appears on bot list page. It can help you identify the purpose of your bot.

Maximum 2000 characters.

AWS | Lex | Lex | X |

Amazon Lex X

Bots

Bot templates New

Networks of bots New

Test workbench New

Test sets

Test results

Related resources

Return to the V1 console

IAM permissions Info

IAM roles are used to access other services on your behalf.

Runtime role
Choose a role that defines permissions for your bot. To create a custom role, use the IAM console.

Create a role with basic Amazon Lex permissions.

Use an existing role.

(i) Creating a role takes a few minutes. Don't delete the role or edit the trust or permissions policies in this role until we've finished creating it.

New role
Amazon Lex creates a runtime role with permission to upload to Amazon CloudWatch Logs.

Bot error logging Info

Debug unexpected issues on Lex bots.

Error logs

Enabled

Disabled

[Learn more about error logs](#) i

[Go to error logs](#) i

Children's Online Privacy Protection Act (COPPA) Info

Is use of your bot subject to the [Children's Online Privacy Protection Act \(COPPA\)](#) i?

Yes

No

Step 1
Configure bot settings

Step 2
Add languages

Add language to bot Info

▼ Language: English (US)

Select language
English (US)

Description - optional

Voice interaction
The text-to-speech voice that your bot uses to interact with users.
Danielle

Voice sample
Hello, my name is Danielle. Let me know how I can assist you. Play

Intent classification confidence score threshold
0.40
Min: 0.00, max: 1.00.

Cancel Add another language Done

- Once the bot is created, add a new **intent**. Name it something like OrderStatus.

Amazon Lex X

Successfully created bot: CustomerSupportBot

Lex > Bots > Bot: CustomerSupportBot > Versions > Version: Draft > All languages > Language: English (US) > Intents > Intent: NewIntent

Draft version English (US) Not built

Intent: NewIntent Info

An intent represents an action that fulfills a user's request. Intents can have arguments called slots that represent variable information.

We've added an intent to get you started. X

Conversation flow Info

Intent details Info

Intent name
OrderStatus
Maximum 100 characters. Valid characters: A-Z, a-z, 0-9, -, _

Intent and utterance generation description
Describe the purpose of your intent. This will also be used when generating utterances for your intent.

Maximum 2000 characters.

ID: CL6RGHGIJZ

- In the intent configuration, add several **sample utterances** (training phrases) that a user might say, such as:
 - Where is my order?
 - Can I get my order status?
 - Check on my package

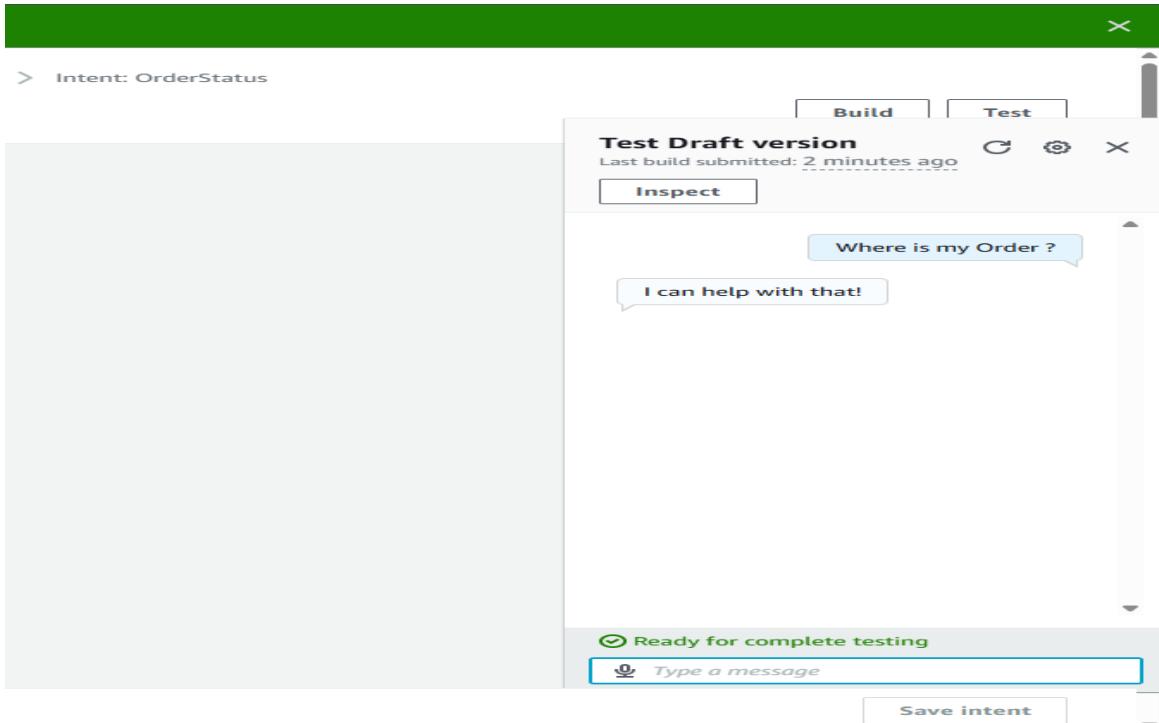
6. Scroll down to the "Closing response" section and add a message that the bot will reply with, like I can help with that! Please provide your order ID to check its status.

7. Save the intent, and then click the "Build" button to train the bot's NLP model.

8. After the build is complete, use the "Test" window within the Lex console to interact with your bot and verify its responses.



Expected Output: When a user types a phrase like "Where is my order?" into the test chat window, the bot should respond with your predefined closing response, "I can help with that! Please provide your order ID to check its status."



Conclusion: Amazon Lex provides a powerful, fully managed service for building conversational AI interfaces using Natural Language Processing (NLP), allowing for the creation of sophisticated chatbots without deep AI expertise.

Practical 2: Text Extraction from Documents using AWS Textract

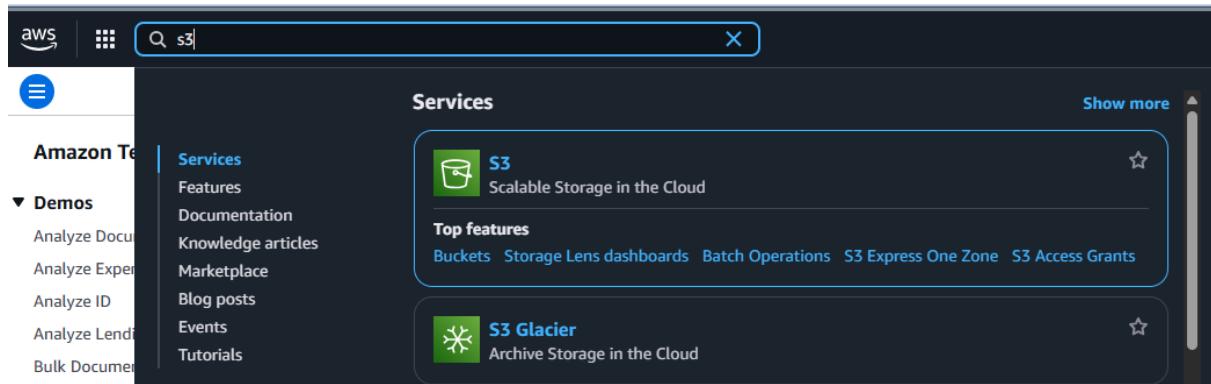
Aim: To use AWS Textract to automatically extract text, forms (key-value pairs), and tables from a scanned document like an invoice or a report.

Tools Used:

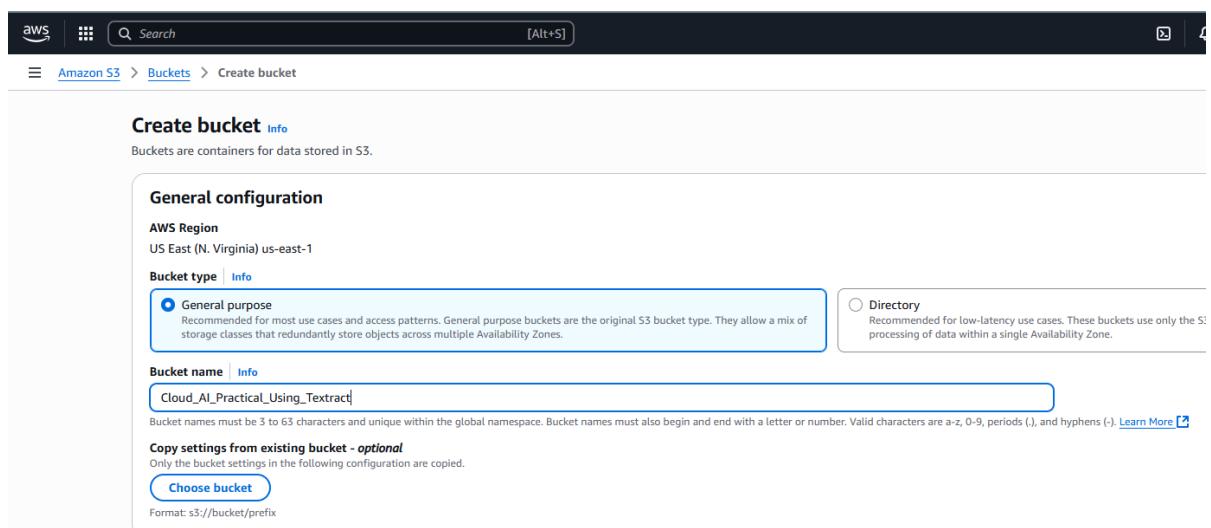
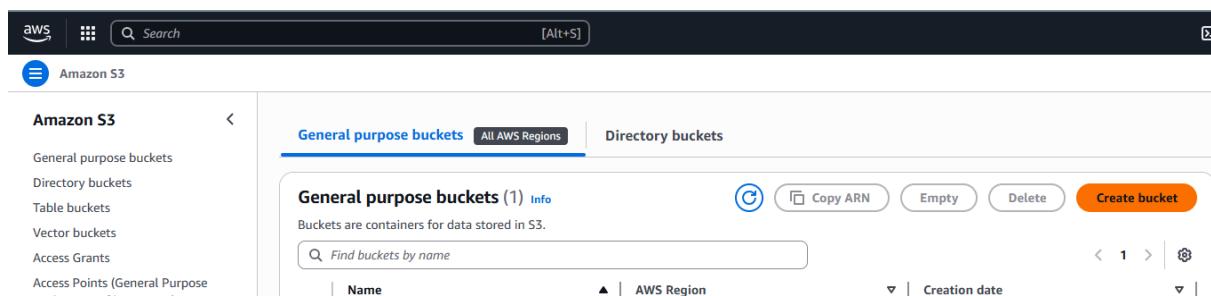
- AWS Management Console
- Amazon S3 (Simple Storage Service)
- AWS Textract

Procedure:

1. Sign in to the AWS Management Console and navigate to the Amazon S3 service.



2. Create a new S3 bucket or use an existing one. Upload a sample document (e.g., an invoice in PDF, JPG, or PNG format) to your bucket.



3. Navigate to the AWS Textract service.

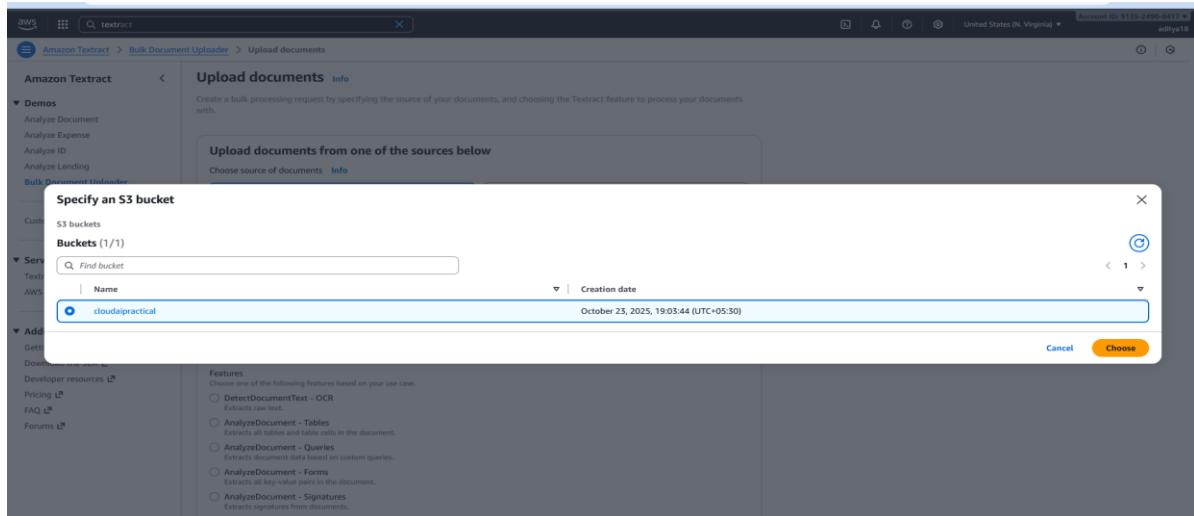
The screenshot shows the AWS Services page. At the top, there's a search bar with the text "texttract". Below it, the "Services" section is displayed, featuring three cards: "Amazon Textract" (blue icon), "Amazon Augmented AI" (green icon), and "AWS End User Messaging" (red icon). Each card has a brief description and a star icon for favoriting.

4. In the Textract console, select "Analyze Document" from the left-hand menu.

The screenshot shows the Amazon Textract service page. The left sidebar includes options like "Analyze Document" under "Demos". The main content area features the heading "Amazon Textract" and the sub-headline "Automatically extract handwriting, text or data from any document using machine learning". A detailed description of the service is provided below.

5. Choose the option to select a document from your Amazon S3 bucket and **browse** to the file you just uploaded.

The screenshot shows the "Bulk Document Uploader" page within the Amazon Textract service. The left sidebar lists "Analyze Document" under "Demos" and "Bulk Document Uploader" under "Additional resources". The main content area is titled "Upload documents" and includes sections for "Upload documents from one of the sources below" (with options for "Import documents from S3 bucket" and "Upload documents from your computer"), "S3 URL" (with a text input field and "View" and "Browse S3" buttons), and "Choose a feature" (with a list of seven feature options: DetectDocumentText - OCR, AnalyzeDocument - Tables, AnalyzeDocument - Queries, AnalyzeDocument - Forms, AnalyzeDocument - Signatures, and AnalyzeDocument - Layout).



6. Textract will automatically process the document. On the results screen, you can view the extracted **raw text**, **forms** (key-value pairs like "Invoice Number": "INV-101"), and **tables**.

Submitted documents (1/2) [Info](#)

Name	Status	Upload date	Type	Feature	Size
Cloud AI Practical Journal	Ready for Download	October 23, 2025, 19:10 (UTC+5:30)	pdf	DetectDocumentText - OCR	434.42 KB

Recent download history

- Cloud AI Practical Journal (1,108 KB • Done)

Full download history

During production use Textract's APIs.

Expected Output: The service will display the document's contents, correctly identifying and separating fields like "Invoice Number," "Due Date," "Total Amount," and any line items within a table.

Conclusion: AWS Textract is more than just OCR; it's an intelligent document processing service that understands document structure, enabling the automation of data entry and document analysis workflows.

Practical 3: Text Sentiment Analysis using AWS Comprehend

Aim

To perform sentiment and entity analysis on text data using **AWS Comprehend**, a cloud-based Natural Language Processing (NLP) service.

Objective

To understand how AWS Comprehend analyzes unstructured text and extracts useful information such as sentiment, entities, and key phrases through pre-trained AI models.

Tools / Services Used

- AWS Management Console
- AWS Comprehend Service
- Sample Text Dataset

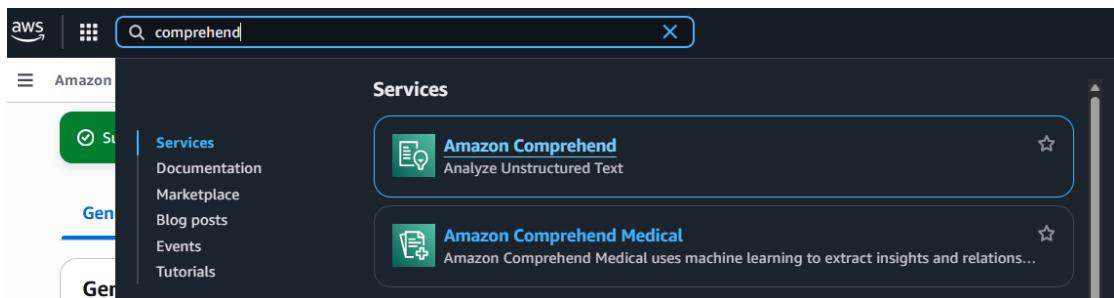
Detailed Procedure

1. Login to AWS Console:

- Open <https://aws.amazon.com/console> and log in.

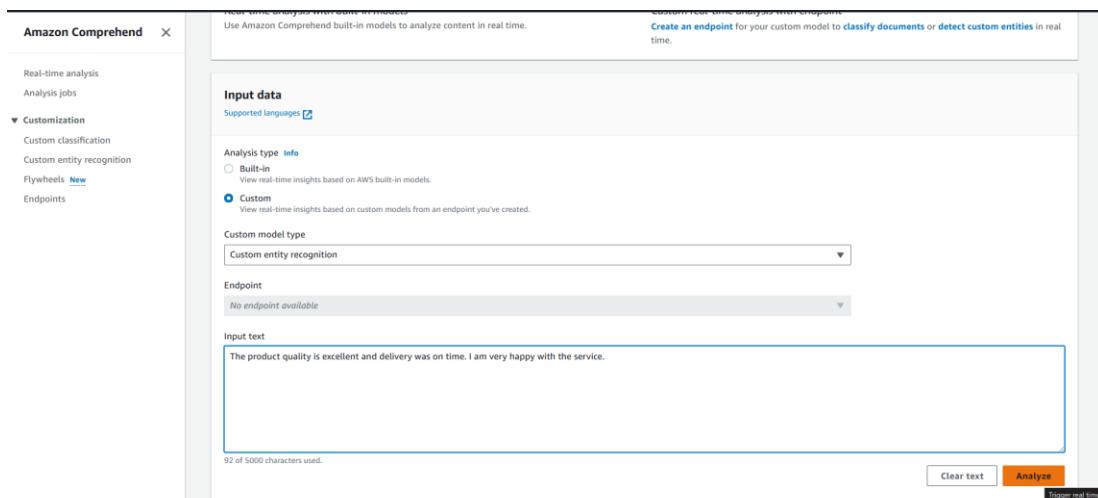
2. Open the Comprehend Service:

- In the search bar, type “**Comprehend**” and open the service dashboard.



3. Select the Analysis Type:

- Choose “**Sentiment Analysis**” to detect emotions in text.
- Optionally, explore “**Entity Recognition**” and “**Key Phrases Extraction**.”



4. Input Sample Text:

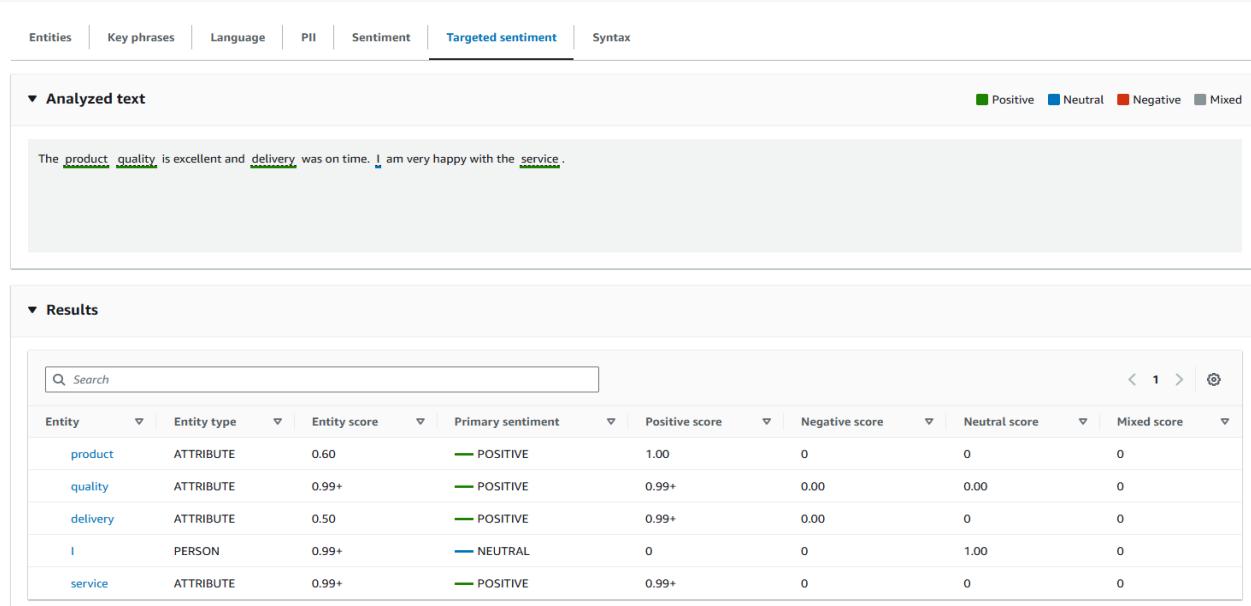
- Enter a paragraph such as:
“The product quality is excellent and delivery was on time. I am very happy with the service.”

5. Run Analysis:

- Click on “Analyze” or “Run” to begin text processing.
- AWS Comprehend uses its NLP models to evaluate sentiment and entities.

View Results:

- The output displays:
 - Overall Sentiment: *Positive*
 - Sentiment Scores: (Positive = 0.98, Negative = 0.01, Neutral = 0.01)
 - Key Entities: *Product, Delivery, Service*
 - Key Phrases: *Excellent quality, On time, Happy with service*



Expected Output

AWS Comprehend classifies text into sentiment categories and highlights entities and key phrases with corresponding confidence scores.

Conclusion

AWS Comprehend enables automatic sentiment and entity analysis without requiring custom NLP development. It can be applied in industries like e-commerce, customer service, and marketing to derive actionable insights from text data.

Practical 4: Image Object and Face Detection using AWS Rekognition

Aim

To analyze and detect objects, scenes, and faces in images using **AWS Rekognition**, an AI-based image and video analysis service.

Objective

To understand how cloud-based computer vision systems perform object detection and facial recognition using pre-trained AI models.

Tools / Services Used

- AWS Management Console
- AWS Rekognition Service

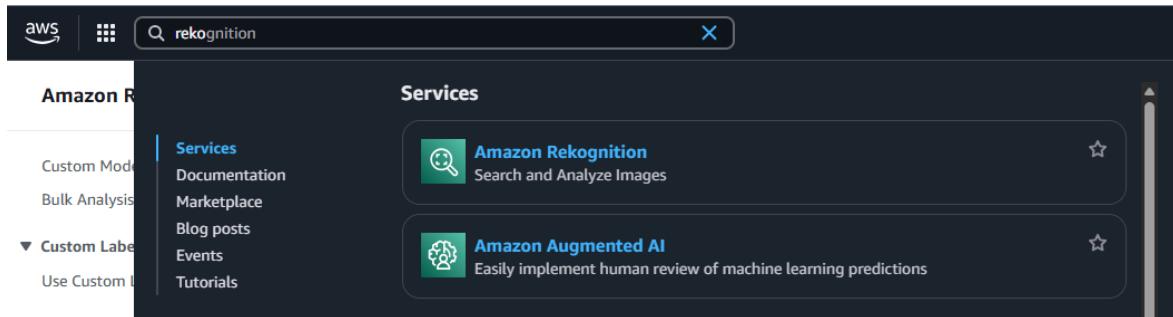
Detailed Procedure (Step-by-Step)

1. Login to AWS Console:

- Go to <https://aws.amazon.com/console> and sign in.

2. Open Rekognition Service:

- Type “Rekognition” in the search bar and select the service.



3. Select the Use Case:

- From the dashboard, choose “Image Properties.”

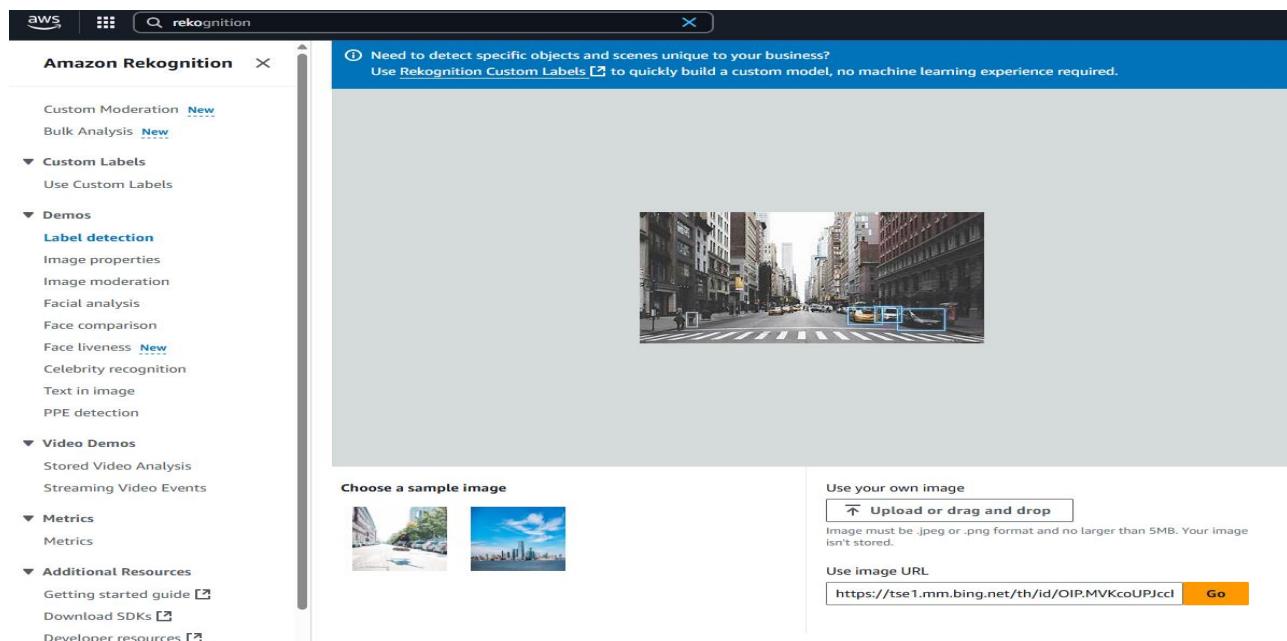
A screenshot of the "Image Properties" page in the Amazon Rekognition service. The main area shows a large image of a city skyline at night, reflected in water. On the left, there's a sidebar with links to "Custom Moderation", "Custom Labels", "Demos", "Metrics", and "Additional Resources". The main content area has tabs for "Image Properties" and "Results". Under "Image Properties", there are sections for "Dominant Colors" (with a list of hex codes and percentages) and "Image Quality" (with metrics for Brightness, Sharpness, and Contrast). Under "Results", there are sections for "Full Image Properties", "Foreground Properties", and "Background Properties". At the bottom, there are options to "Choose a sample image" or "Use your own image" and a "Request" button.

4. Upload or Select Image:

- Upload a local image
- Use a clear image containing people or objects (e.g., a group photo or street scene).

5. Run Object and Scene Detection:

- Click “**Detect Labels.**”
- Rekognition processes the image and lists detected items (e.g., “Person,” “Car,” “Building,” “Tree”).



6. View Confidence Scores:

The screenshot shows the results page of the Amazon Rekognition console. At the top, a banner encourages using Rekognition Custom Labels to quickly build a custom model. Below the banner, there is a search bar labeled "Search all labels" and a link to "Download full list". A search input field contains "Enter a label name". The main content area is titled "Results" and lists the following confidence scores:

Label	Confidence (%)
City	100 %
Road	100 %
Street	100 %
Urban	100 %
Neighborhood	99.9 %
Metropolis	99.9 %

A "Show more" button is located at the bottom of the list.

- Each label is displayed with a confidence percentage (e.g., Person – 99.8%, Car – 97.2%).

Conclusion

AWS Rekognition demonstrates the power of cloud-based computer vision systems in recognizing and classifying image content. It simplifies the development of intelligent visual applications for industries like security, retail, and automation.